

Deploy a Web Application Using a Kubernetes Deployment and Service

Table of Contents

1. **Introduction**
 2. **Problem Statement**
 3. **Prerequisites**
 4. **Setup Instructions**
 - **Step 1: Create a Deployment for NGINX**
 - **Step 2: Create a Service to Expose NGINX**
 - **Step 3: Accessing the NGINX Application**
 5. **References**
-

Introduction

This guide is designed to help users set up and deploy web applications in a local Kubernetes environment using Minikube on a Windows system. Minikube is an open-source tool that runs a single-node Kubernetes cluster on your local machine, providing an easy-to-use environment for Kubernetes learning, development, and testing.

By following this guide, you will learn:

- How to set up a Kubernetes cluster locally using Minikube.
- How to deploy web applications using Kubernetes Deployment and manage their lifecycle.
- How to expose your web applications to internal or external traffic using Kubernetes Service.

Problem Statement

In this lab, we will deploy an NGINX-based web application on a local Kubernetes cluster using **Minikube**. We will create a Kubernetes Deployment to manage the application's lifecycle and a Service to expose it.

Prerequisites

Completion of all previous lab guides (up to Lab Guide-01) is required before proceeding with Lab Guide-02.

- Minikube is running on your Windows system.
- `kubect1` is installed and configured to interact with your Minikube cluster.

Setup Instructions

Step 1: Create a Deployment for NGINX

We will start by deploying NGINX using a Kubernetes Deployment. The Deployment ensures that NGINX is running in a stable state with three replicas.

1. **Create the YAML for Deployment**

Create a new file called `nginx-deployment.yaml` in your working directory with the following content:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Key points:

- **replicas:** 3 instances of NGINX will be deployed.
- **image:** NGINX is pulled from the official Docker image (`nginx:latest`).
- **ports:** Exposes port 80 inside the container.

2. Deploy NGINX to Kubernetes

Apply the `nginx-deployment.yaml` file by running the following command:

```
kubectl apply -f nginx-deployment.yaml
```

3. Verify the Deployment

Check the status of the deployment and ensure the pods are running:

```
kubectl get deployments
kubectl get pods
```

```
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    0/3     3            0           18s
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-54b9c68f67-rpx58   1/1     Running   0          38s
nginx-deployment-54b9c68f67-tm8dh   1/1     Running   0          38s
nginx-deployment-54b9c68f67-zwfsh   1/1     Running   0          38s
```

You should see three pods running for the `nginx-deployment`.

Step 2: Create a Service to Expose NGINX

Next, we will create a Service to expose the NGINX application inside the cluster. A **ClusterIP** service will be used to make the application accessible to other pods within the cluster.

1. Create the YAML for Service

Create a new file called `nginx-service.yaml` with the following content:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - port: 80
      targetPort: 80
```

Key points:

- **type: ClusterIP:** Exposes the service within the cluster (internal access).
- **selector:** Targets the pods labeled with `app: nginx`.
- **ports:** Exposes port 80 both internally and in the container.

2. Deploy the Service

Apply the `nginx-service.yaml` by running the following command:

```
kubectl apply -f nginx-service.yaml
```

3. Verify the Service

Check the service status and ensure it's running:

```
kubectl get services
```

```
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl apply -f nginx-service.yaml
service/nginx-service created
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl get services
NAME                TYPE        CLUSTER-IP      EXTERNAL-IP  PORT(S)    AGE
kubernetes          ClusterIP   10.96.0.1       <none>       443/TCP    22h
nginx-service       ClusterIP   10.107.163.89   <none>       80/TCP     6s
```

You should see the **nginx-service** listed.

Step 3: Accessing the NGINX Application

Since this service is internal to the cluster, you can use **kubectl** to forward the local port to access NGINX from your browser.

1. Port-forward the Service

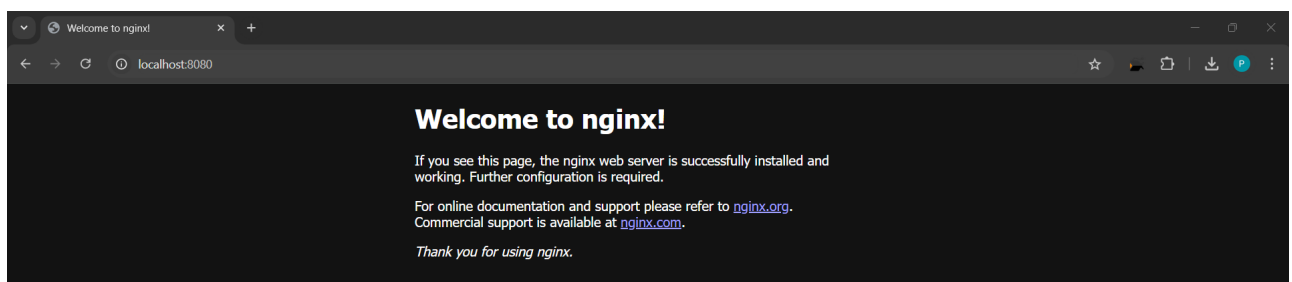
Run the following command to forward a local port (e.g., 8080) to the NGINX pod:

```
kubectl port-forward service/nginx-service 8080:80
```

```
(base) PS D:\GuideLabs\Guided_Labs\Kubernetes\k8s_Example> kubectl port-forward service/nginx-service 8080:80
Forwarding from 127.0.0.1:8080 -> 80
Handling connection for 8080
```

2. Open NGINX in Your Browser

Open your browser and navigate to **http://localhost:8080**. You should see the default NGINX welcome page.



References

- [Kubernetes Deployment Documentation](#)
- [Kubernetes Service Documentation](#)